



Hochschule für Angewandte Wissenschaften Hamburg
Hamburg University of Applied Sciences

Evaluation of parallelism techniques on embedded multi-core platforms

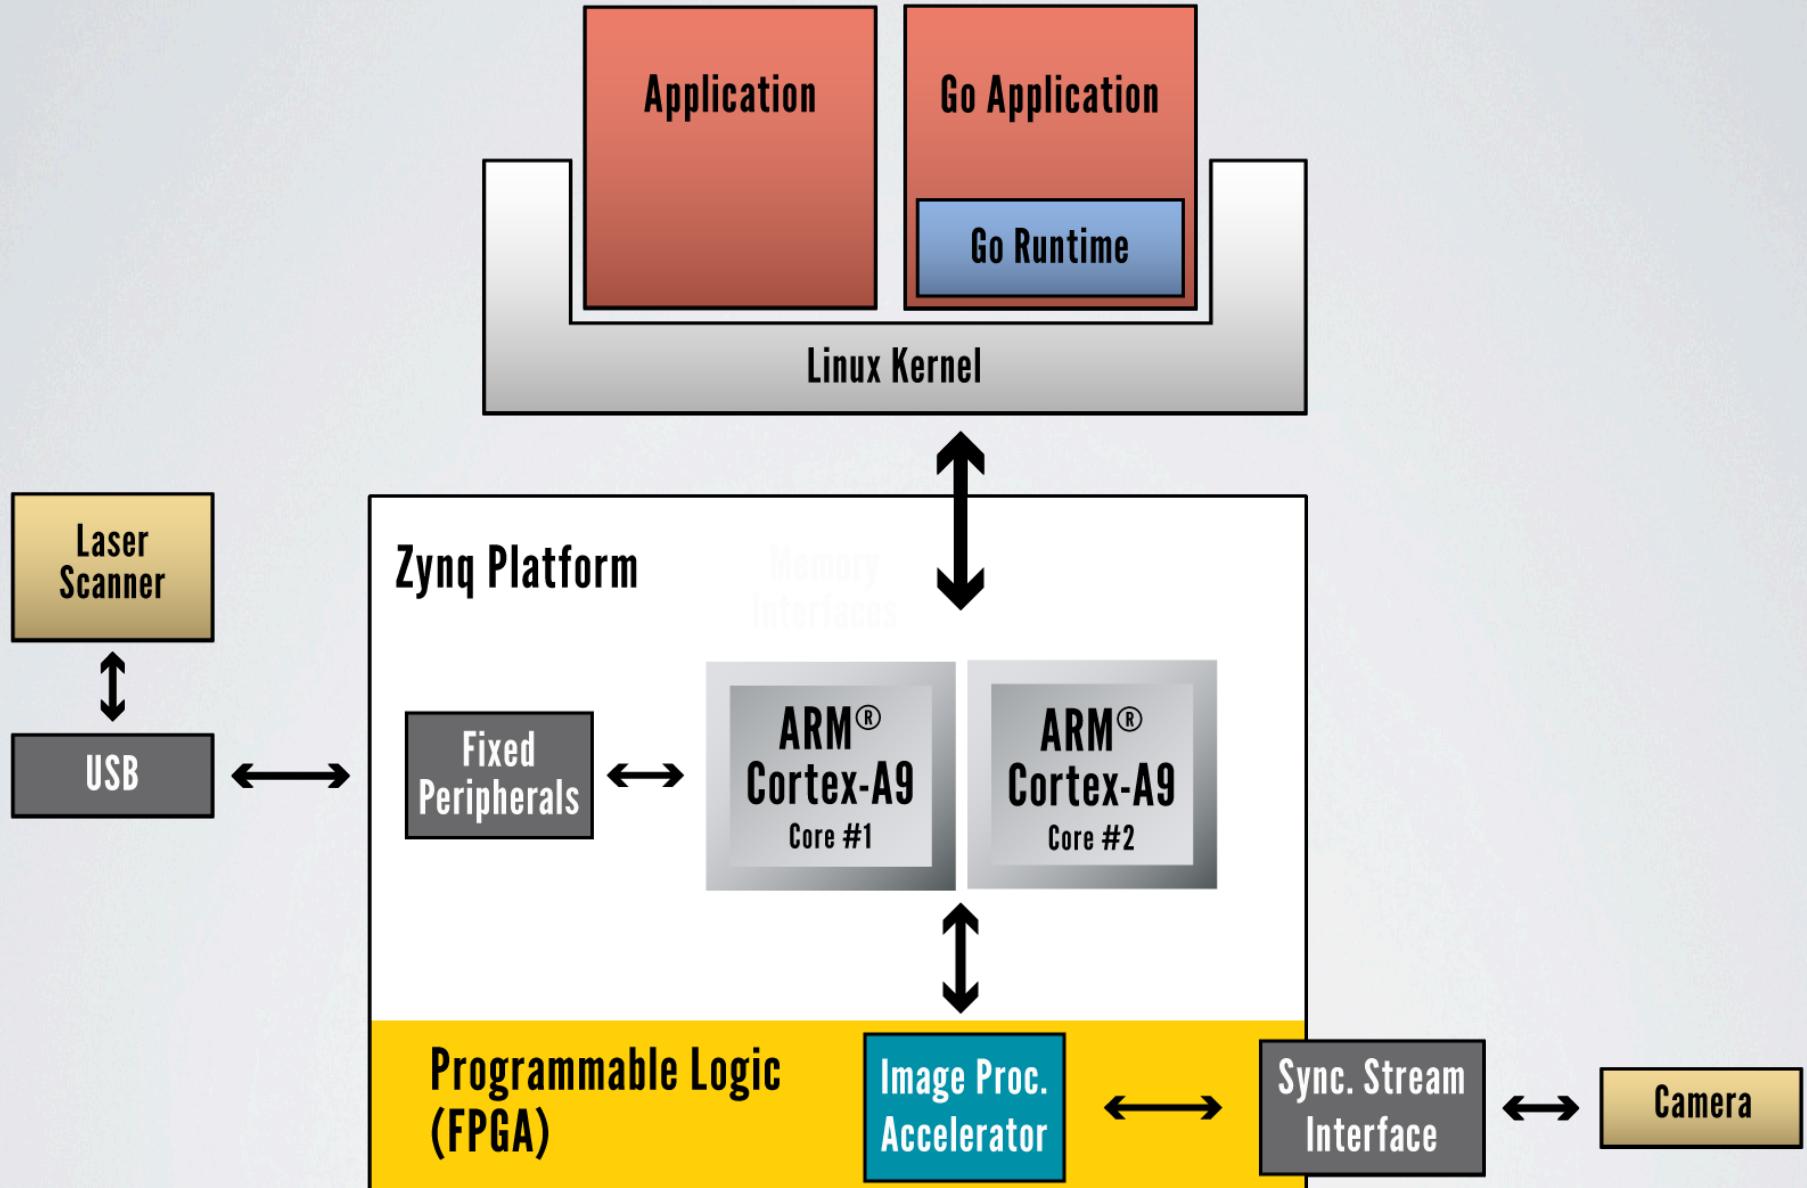
Lucas Jenß

Mentor: Prof. Dr. B. Schwarz

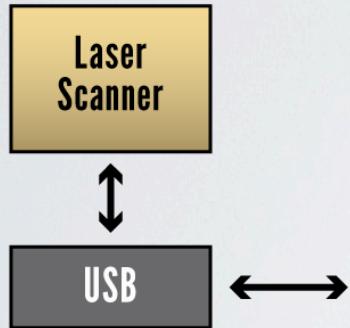
Outline

1. Development Platform
2. Parallelism Techniques
 - Pipeline Pattern
 - Single Instruction – Multiple Data (SIMD)
3. Google Go for parallelism

The Platform



Obstacle Detection



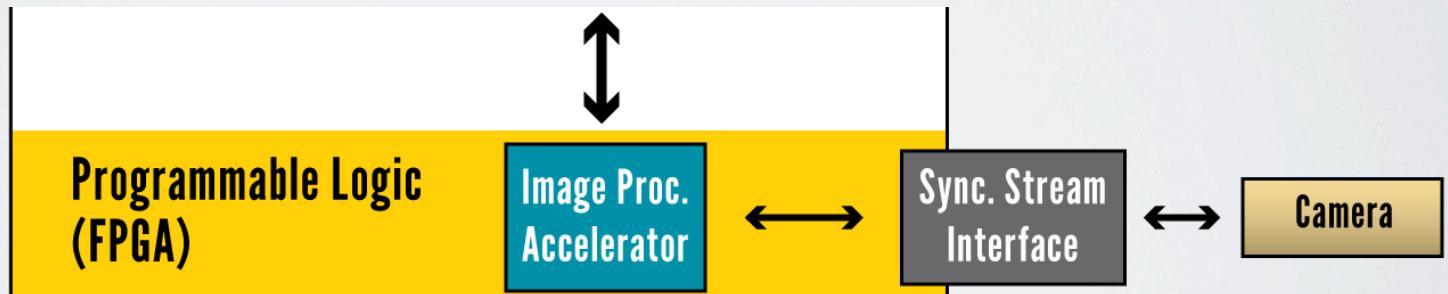
- Using Laser scanner
 - Data rate: ~34KB/s
 - 10Hz scanning freq.
 - 240° in 0.36° steps
- Currently implemented using Java on Android OS

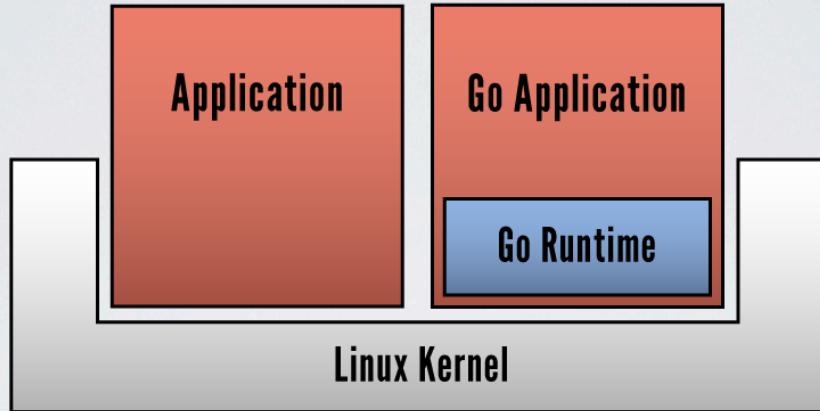


[1]

Lane Guiding Control

- Lane Guiding Control using high FPS camera (60 FPS, 640x480px)
- Data rate: ~17.6MB/s
- Pixel Clock: 27MHz
- Implemented using C and RTL models, on Linux





- Current state: mix of different technologies
 - Java on Android
 - C on Linux parallel to RTL/FPGA
- Goals:
 - technological unification of the platform
 - parallelization of tasks to improve throughput

Outline

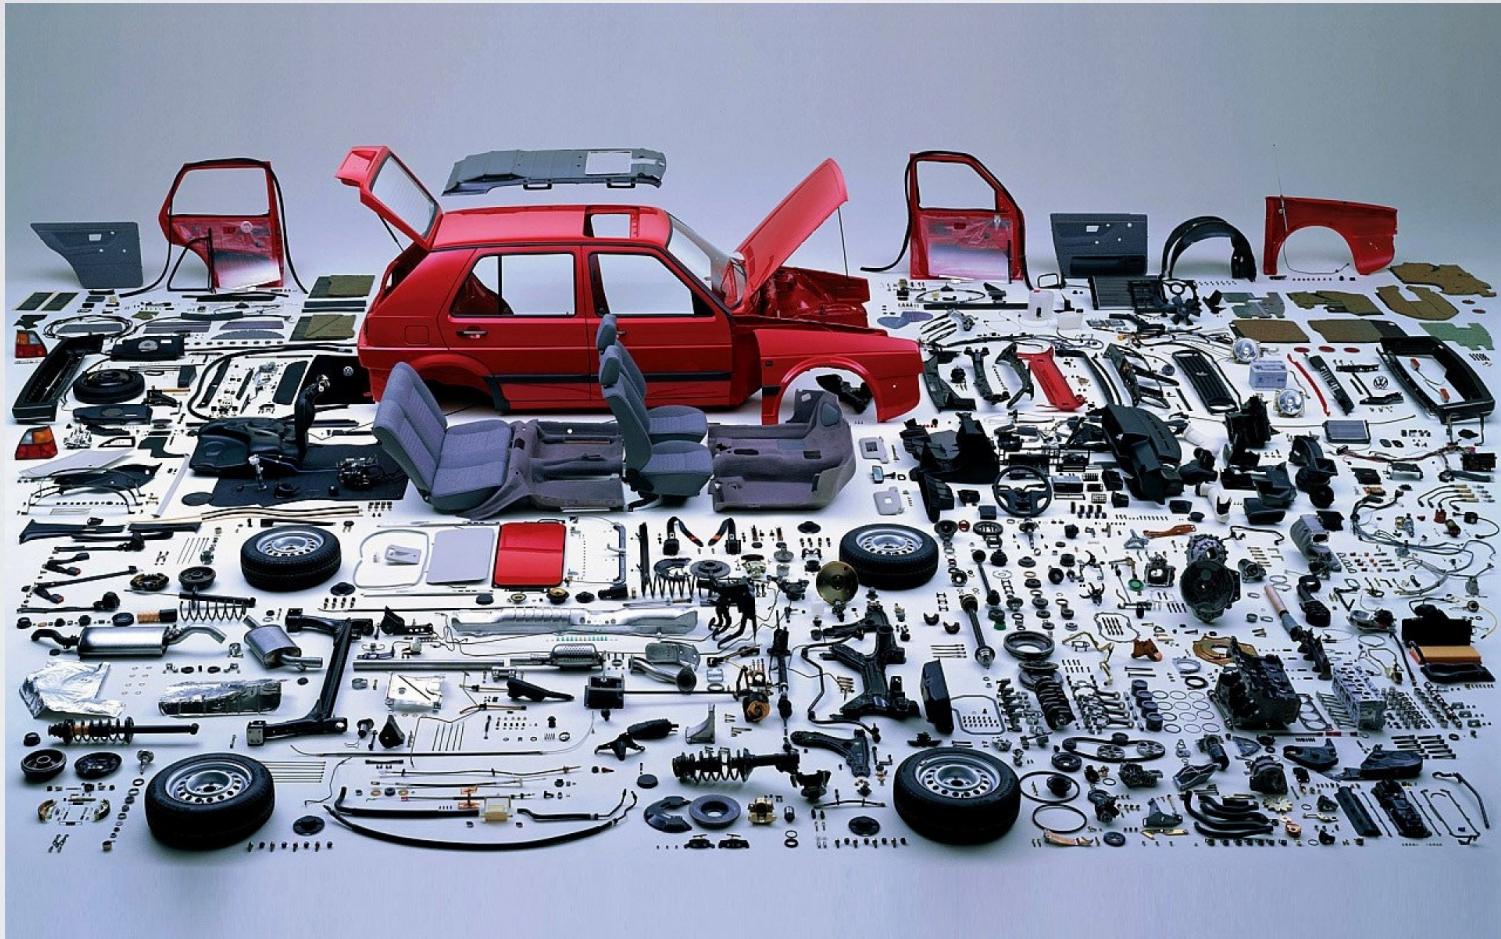
1. Development Platform
2. Parallelism Techniques
 - Pipeline Pattern
 - Single Instruction – Multiple Data (SIMD)
3. Google Go for parallelism

Parallelism Techniques

Pipeline pattern

Pipeline Pattern

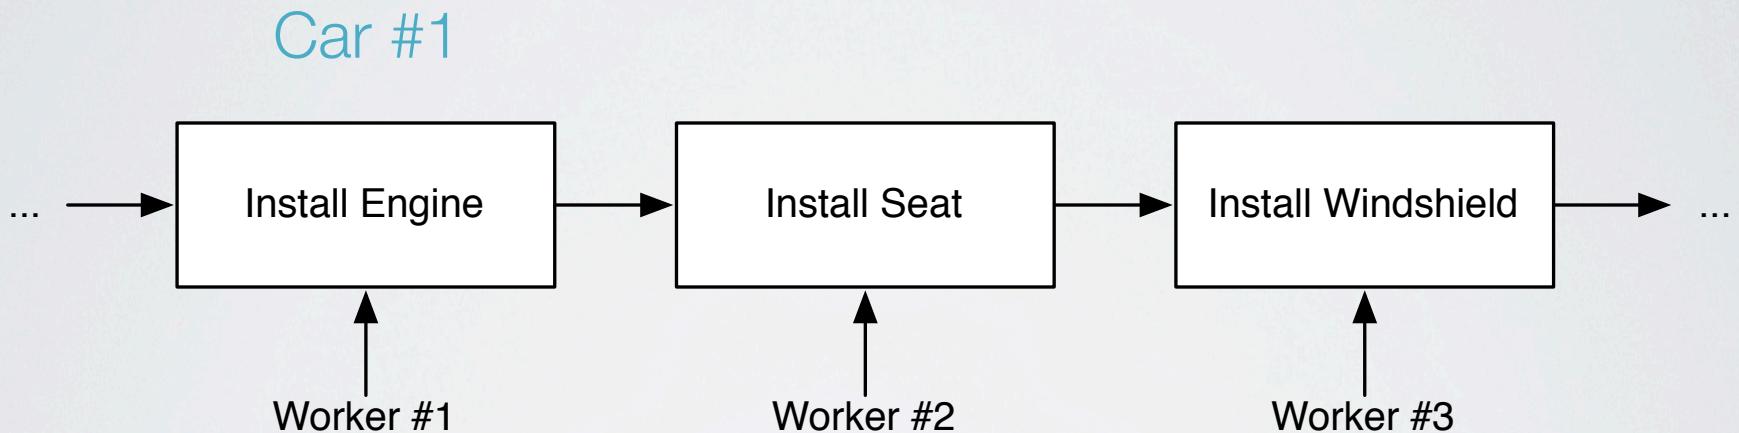
"Assembly line" example



Pipeline Pattern

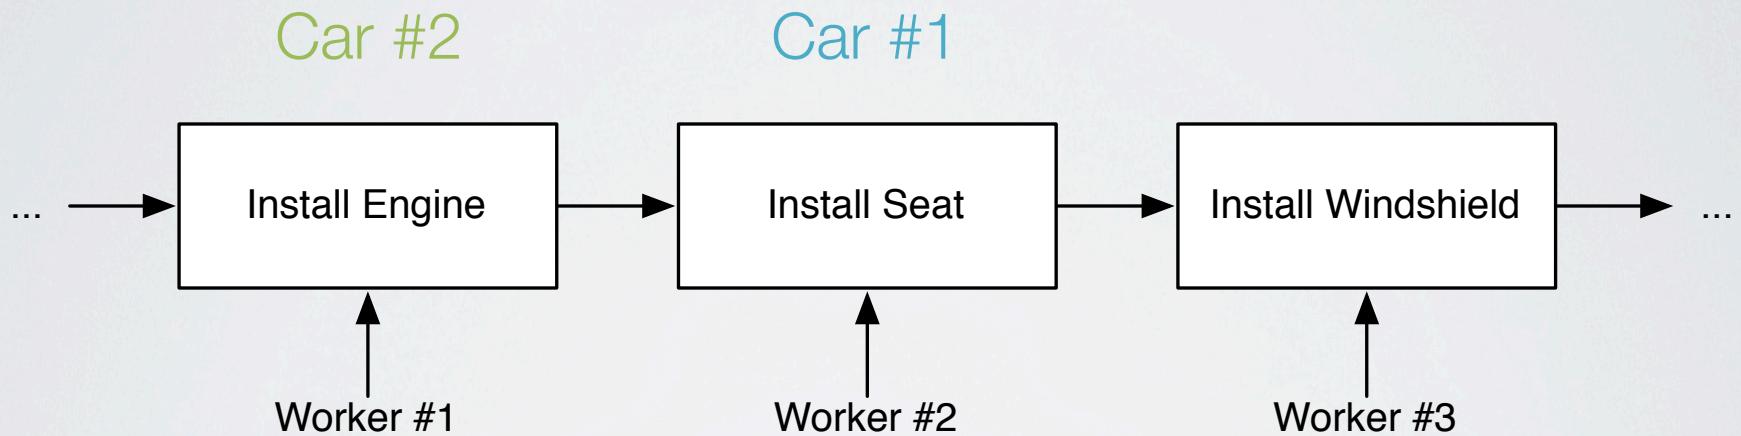
"Assembly line" example

- Analogy: Assembly line (e.g. for a Car)



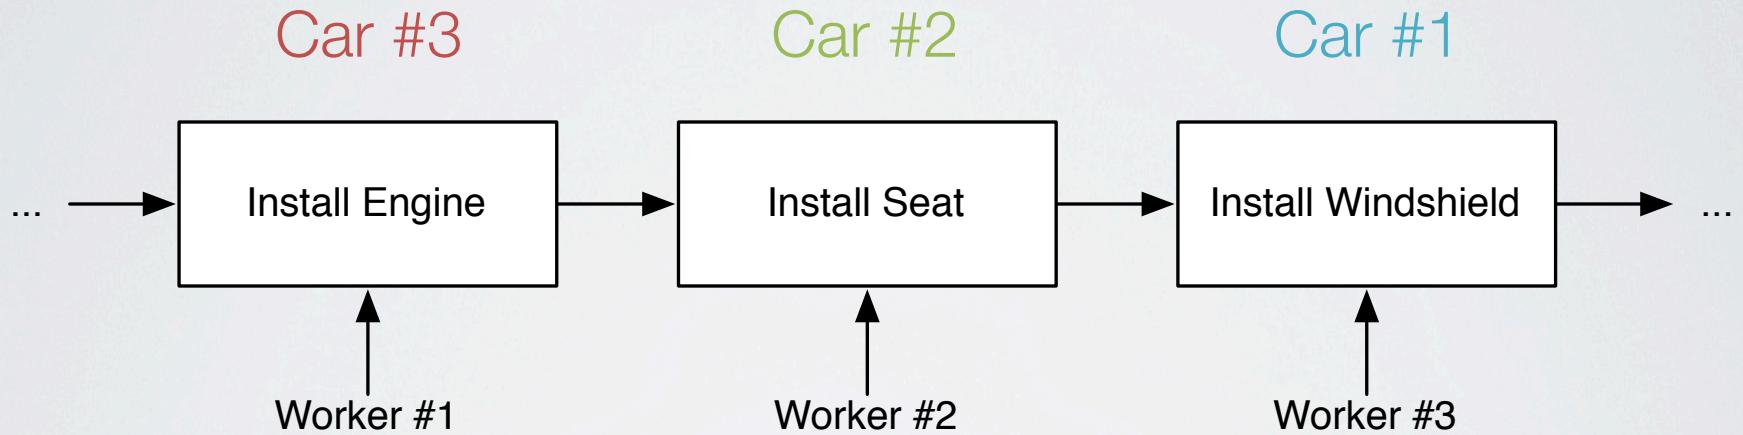
Pipeline Pattern

"Assembly line" example



Pipeline Pattern

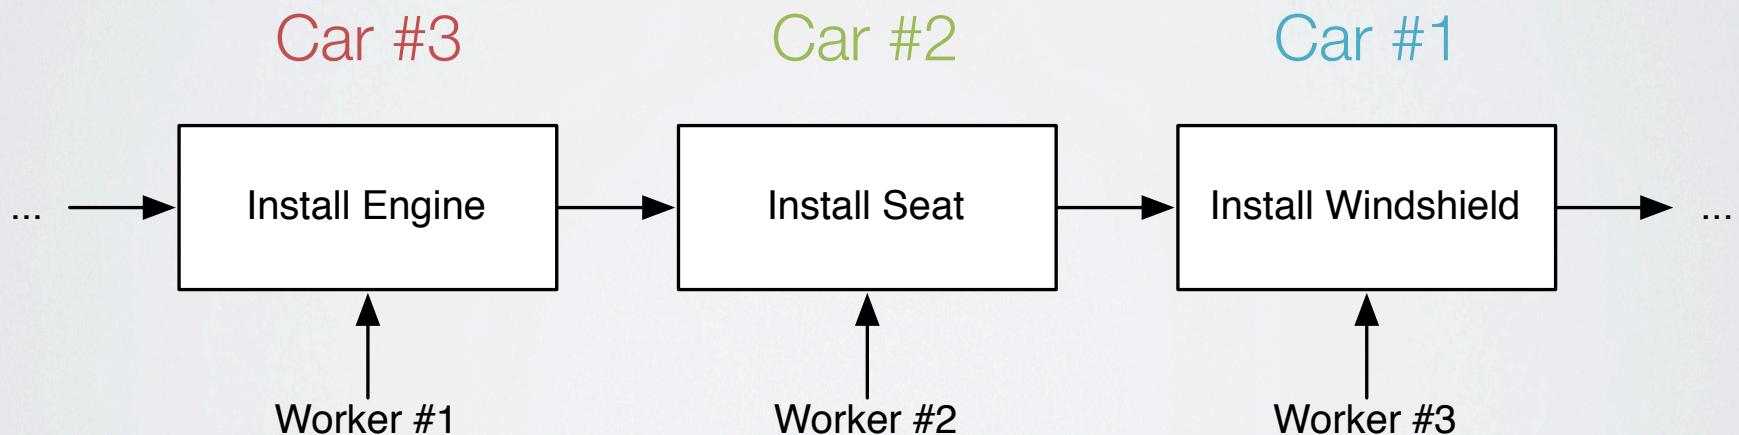
"Assembly line" example



Pipeline Pattern

”Assembly line” example

- All workers
 - are busy at all times
 - work on different tasks on different cars
 - yet all tasks are directly related



Pipeline Pattern

”Assembly line” example

- Assembly line has
 - Workers
 - Cars, or some other items
 - Transport mechanism for the items

Pipeline Pattern

”Assembly line” example

- For the pipeline pattern, this translates to
 - Threads (instead of Workers)
 - Data (instead of Cars)
 - Communication channel
(instead of the transport mechanism)

Pipeline Stage

Pseudo code

```
while(more elements) {  
    receive element from previous stage  
    perform operation on element  
    send element to the next stage  
}
```

Pipeline Stage

Pseudo code

```
while(more elements) {  
    receive element from previous stage  
    perform operation on element  
    send element to the next stage  
}
```

Pipeline Stage

Pseudo code

```
while(more elements) {  
    receive element from previous stage  
    perform operation on element  
    send element to the next stage  
}
```

Pipeline Pattern

- Theoretically:
 - Parallelization of any computing problem
 - Because the pipeline “transforms” linearity into parallelism

Pipeline Pattern

- In practice:
 - Communication between stages is the bottleneck
 - In order to achieve good performance:
 - Communication time \ll Stage operation time
 - Small difference between stage operation times

Outline

1. Development Platform
2. Parallelism Techniques
 - Pipeline Pattern
 - Single Instruction – Multiple Data (SIMD)
3. Google Go for parallelism

Single Instruction – Multiple Data

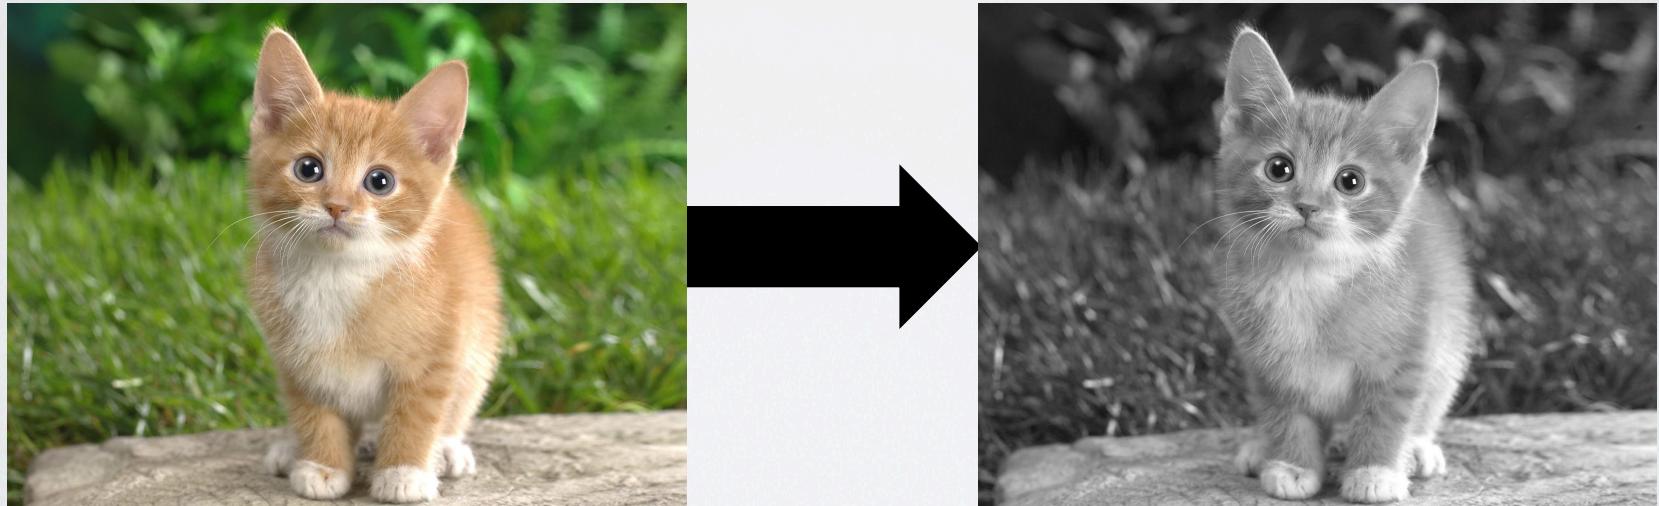
Using the ARM NEON engine

Single Instruction - Multiple Data

- Exploits data-level parallelism
 - Operations repeatedly applied to independent data

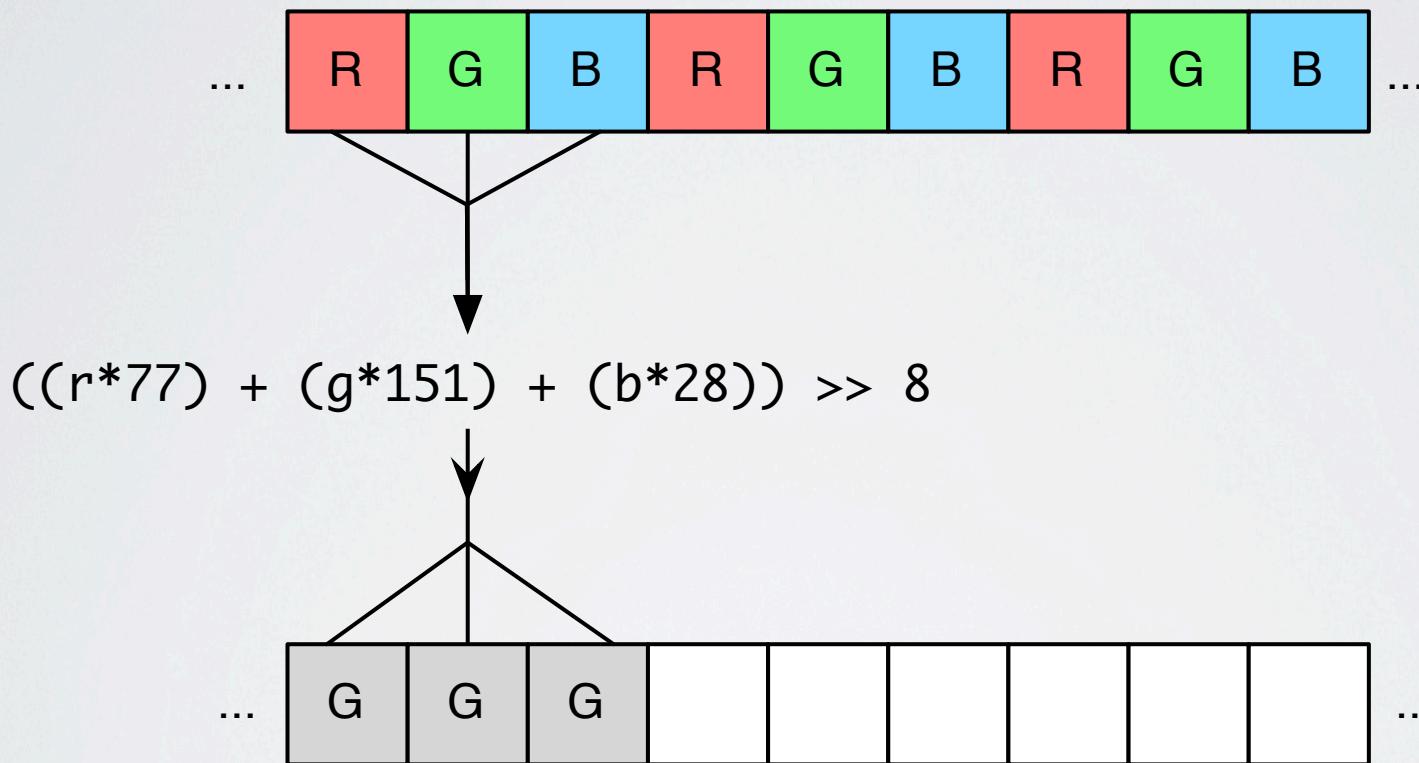
Single Instruction - Multiple Data

- Example: Converting an image to gray scale
 - Weighted average of each pixel's RGB values
 - Sum of products



Single Instruction – Single data

Converting an image to gray scale



Single Instruction – Single data

Image to gray scale – Sum of products

```
void rgb2grey(uint8_t *dest, uint8_t *src, int pixels)
{
    for(int i=0; i < pixels; i++) {
        uint8_t r = *src++;
        uint8_t g = *src++;
        uint8_t b = *src++;

        uint16_t y =
            r * 77 +
            g * 151 +
            b * 28;

        *dest++ = (y >> 8);
    }
}
```

Single Instruction – Single data

Converting an image to gray scale

```
void rgb2grey(uint8_t *dest, uint8_t *src, int pixels)
{
    for(int i=0; i < pixels; i++) {
        uint8_t r = *src++;
        uint8_t g = *src++;
        uint8_t b = *src++;

        uint16_t y =
            r * 77 +
            g * 151 +
            b * 28;

        *dest++ = (y >> 8);
    }
}
```

Single Instruction – Single data

Converting an image to gray scale

```
void rgb2grey(uint8_t *dest, uint8_t *src, int pixels)
{
    for(int i=0; i < pixels; i++) {
        uint8_t r = *src++;
        uint8_t g = *src++;
        uint8_t b = *src++;

        uint16_t y =
            r * 77 +
            g * 151 +
            b * 28;

        *dest++ = (y >> 8);
    }
}
```

Single Instruction – Single data

Converting an image to gray scale

```
void rgb2grey(uint8_t *dest, uint8_t *src, int pixels)
{
    for(int i=0; i < pixels; i++) {
        uint8_t r = *src++;
        uint8_t g = *src++;
        uint8_t b = *src++;

        uint16_t y =
            r * 77 +
            g * 151 +
            b * 28;

        *dest++ = (y >> 8);
    }
}
```

Single Instruction – Multiple data

Converting an image to gray scale

- ARM NEON uses special registers
 - Separate from CPU registers
 - “Q”(uad) and “D”(ouble), 128-bit and 64-bit wide
 - Multiple elements per register
 - e.g. 8 * 8 bit in “D” register

Single Instruction – Multiple data

Converting an image to gray scale

- Steps to use SIMD:
 1. Load data into NEON registers
 2. Apply operations
 3. Write data back to memory

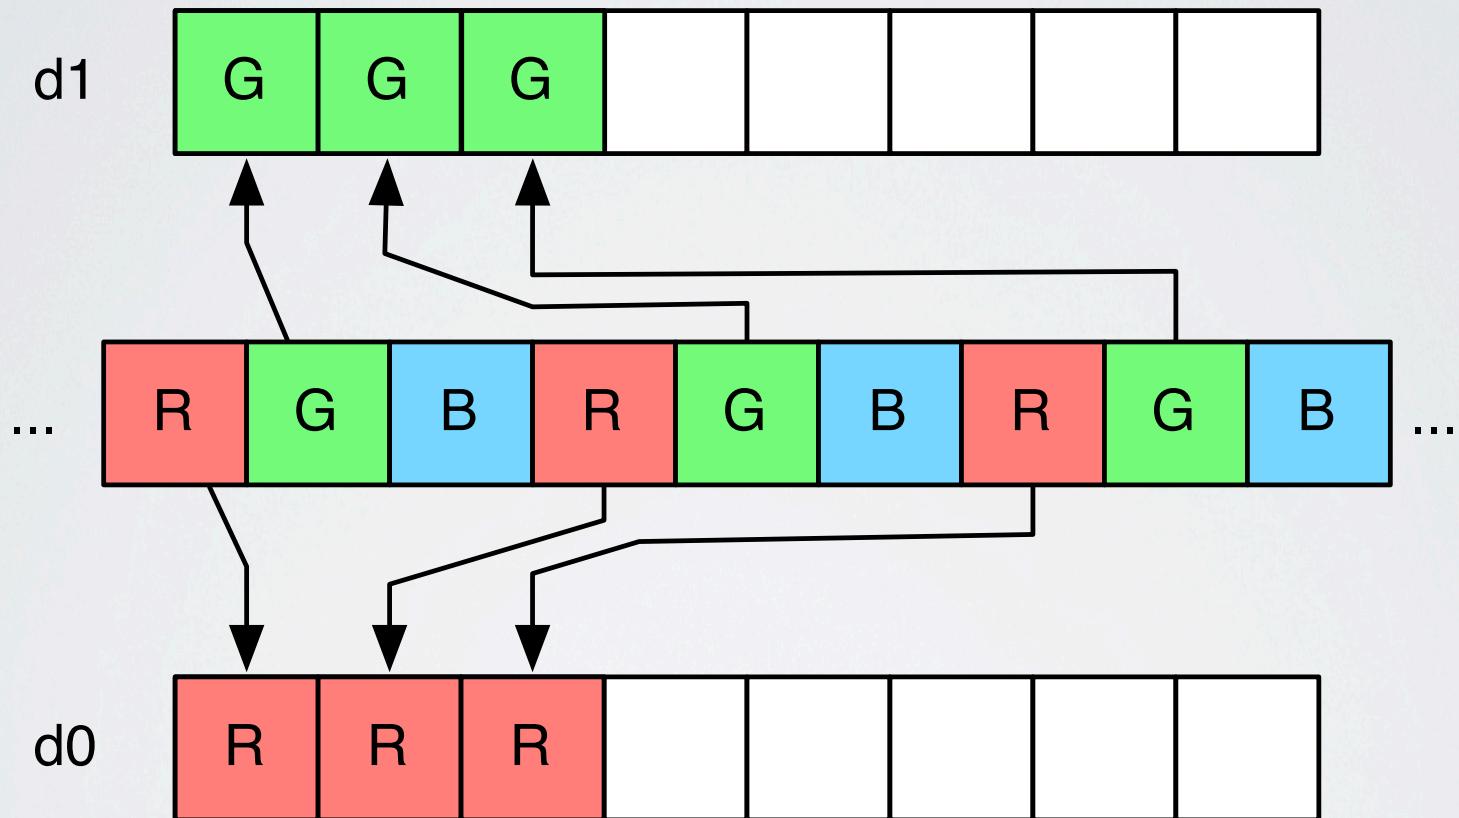
Single Instruction – Multiple data

Converting an image to gray scale

- Steps to use SIMD:
 1. Load data into NEON registers
 2. Apply operations
 3. Write data back to memory

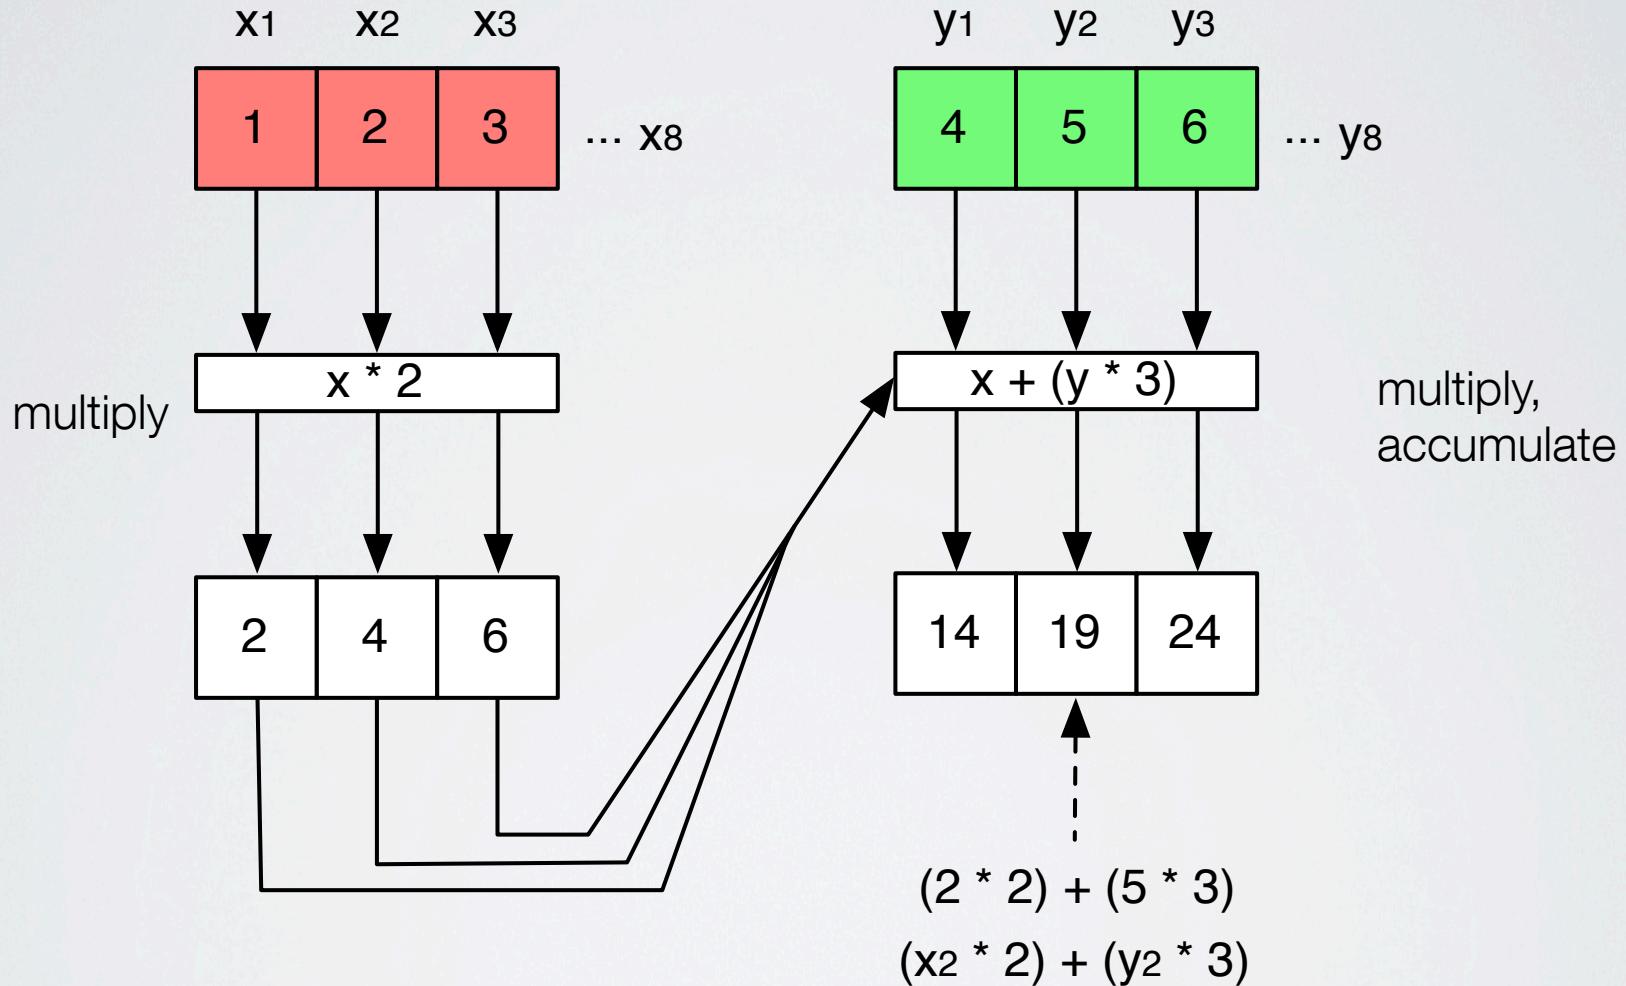
Single Instruction – Multiple data

Load data into NEON “D” registers



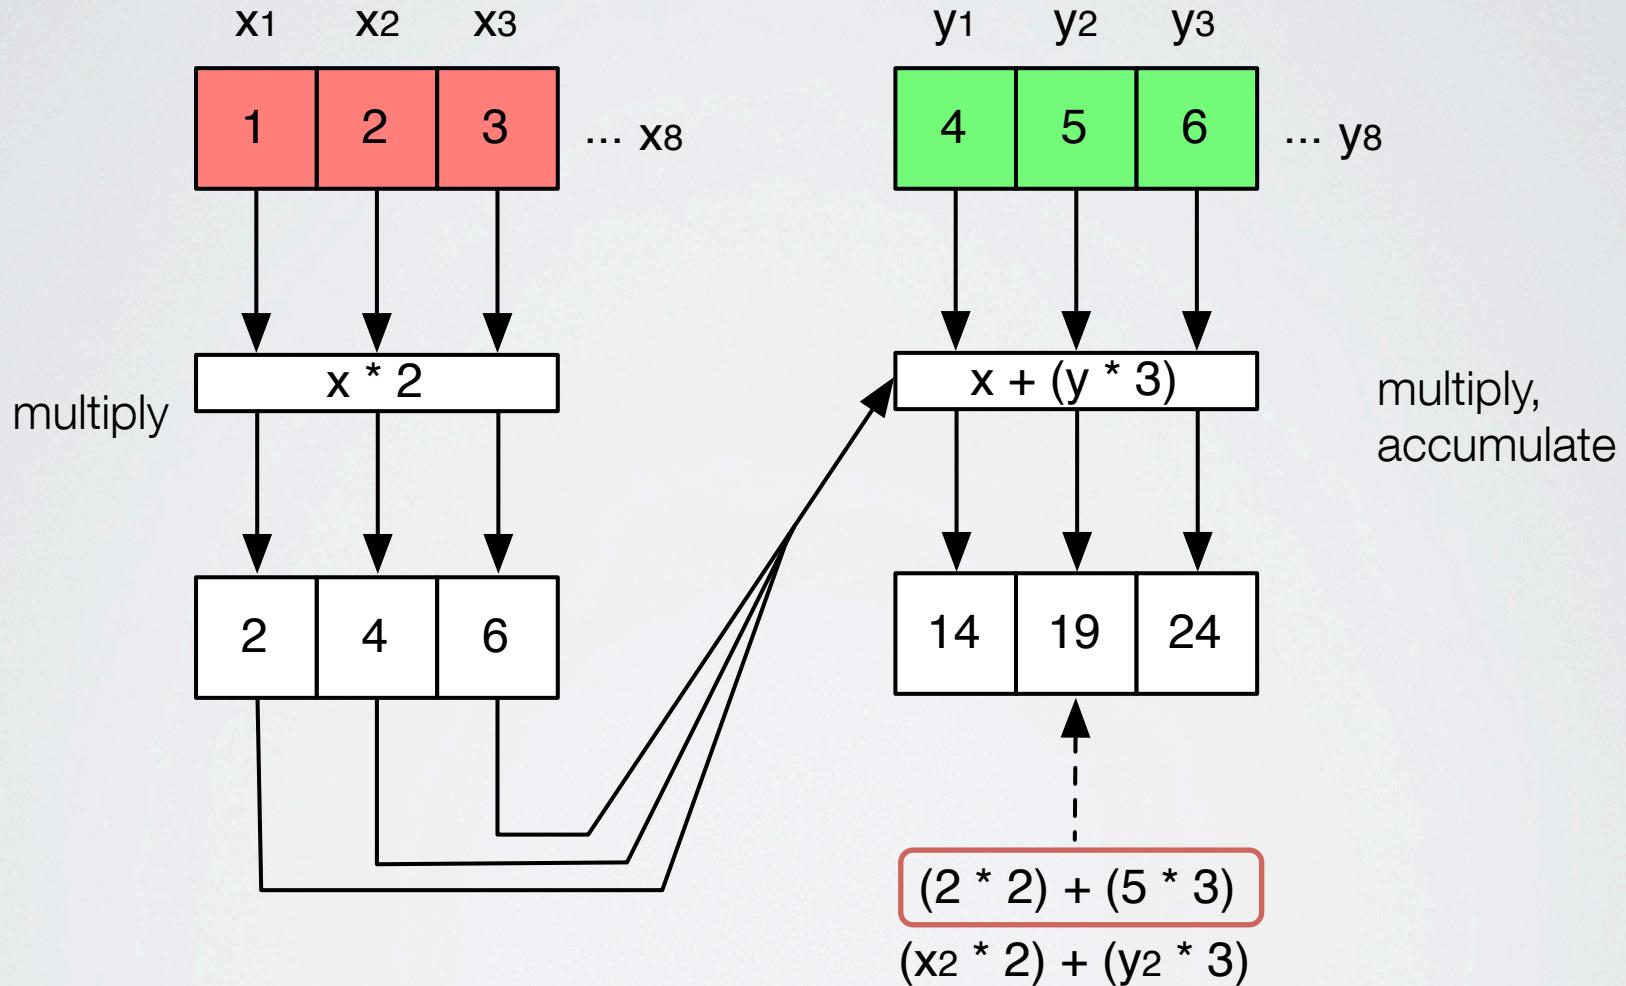
Single Instruction – Multiple data

Apply operations



Single Instruction – Multiple data

Apply operations



Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Load (3 Interleaved)

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;
```

uint8x8x3_t

Type

Elements per Register

}

of Registers

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;
```



Loads 8bit * 8 * 3 = 192 bit = 24 byte!

```
}
```

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Apply operations

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Multiply

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
    Multiply +  
Accumulate  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8); Shift Right  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result); } Write back to memory  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

Single Instruction – Multiple data

Converting to gray scale

```
for(int i=0; i < pixels; i += FACTOR) {  
    uint8x8x3_t rgb = vld3_u8(src);  
    uint16x8_t acc;  
  
    acc = vmull_u8(rgb.val[0], rcoeff);  
    acc = vmlal_u8(acc, rgb.val[1], gcoeff);  
    acc = vmlal_u8(acc, rgb.val[2], bcoeff);  
  
    uint8x8_t result = vshrn_n_u16(acc, 8);  
    vst1_u8(dest, result);  
  
    src += FACTOR;  
    dest += FACTOR;  
}
```

FACTOR = # of elements
processed per iteration (24)

Single Instruction – Multiple data

Performance

- Native C-Version
 - 16 Instructions / Pixel
- SIMD Version:
 - 15 Instructions / 8 Pixels
 - ~2 Instructions / Pixel

Outline

1. Development Platform
2. Parallelism Techniques
 - Pipeline Pattern
 - Single Instruction – Multiple Data (SIMD)
3. Google Go for parallelism

Google Go

for parallel applications

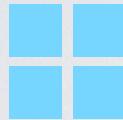
Google Go

- A programming language that is
 - simple
 - compiled
 - built with concurrency in mind

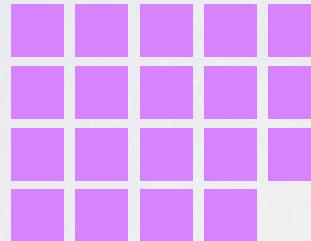
Language Specification

Who's keeping it simple?

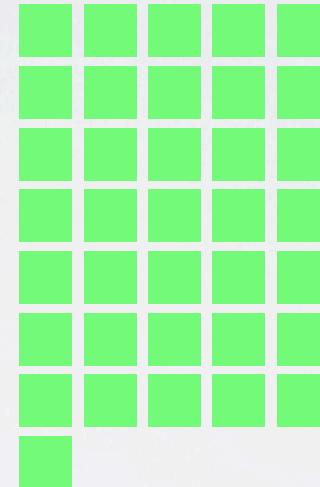
Go



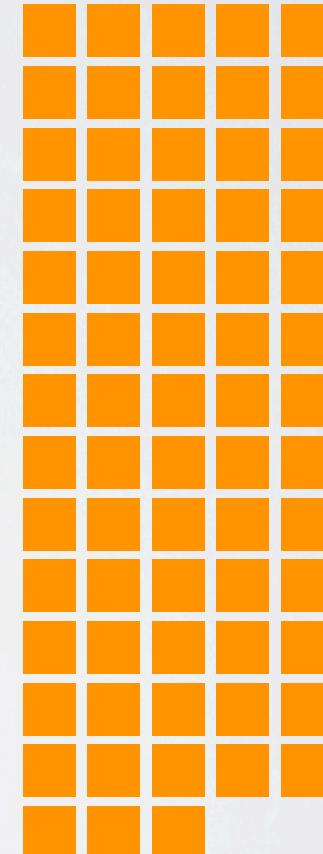
Scala



C++



Java



■ = 10 pages (DIN A4)

Google Go

Concurrency as a language feature

- Go-Routines (**go** keyword)
 - lightweight threads scheduled by the Go runtime
 - concurrently execute *any* function
 - sequential execution:
function(parameter1, parameterN...)

Google Go

Concurrency as a language feature

- Go-Routines (**go** keyword)
 - lightweight threads scheduled by the Go runtime
 - concurrently execute *any* function
 - concurrent execution:
go function(parameter1, parameterN...)

Google Go

Concurrency as a language feature

- Channels
 - first-class value (i.e. handled as a primitive)
 - thread-safe communication mechanism
 - buffered / unbuffered

Google Go

Concurrency as a language feature

- Channels
 - create channel

```
chan := make(chan int)
```

Google Go

Concurrency as a language feature

- Channels
 - send data to channel

```
chan <- 5
```

Google Go

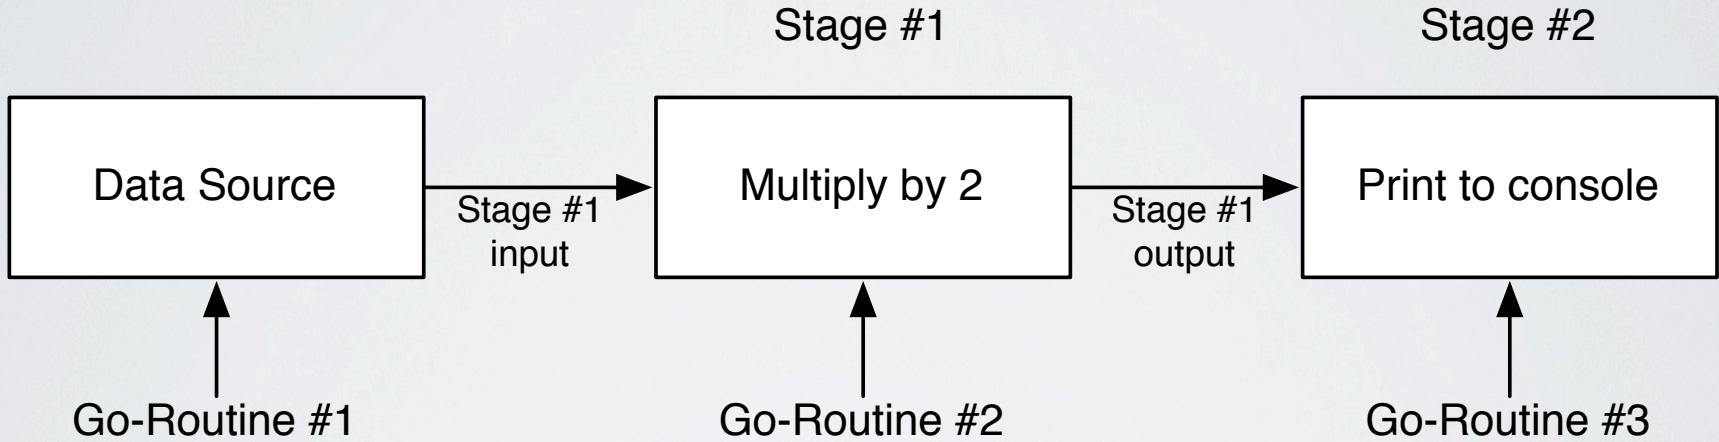
Concurrency as a language feature

- Channels
 - retrieve data from channel

```
value := <- chan
```

Code example

Pipeline pattern in Go!



Code example

Pipeline pattern in Go!

```
func main() {
    stage1_in := make(chan int)
    stage1_out := make(chan int)

    go func() {
        for _, element := range [4]int{1, 2, 3, 4} {
            stage1_in <- element
        }
    }()
}

go stage1(stage1_in, stage1_out)
go stage2(stage1_out, nil)
}
```

Code example

Pipeline pattern in Go!

```
func main() {
    stage1_in := make(chan int)
    stage1_out := make(chan int)

    go func() {
        for _, element := range [4]int{1, 2, 3, 4} {
            stage1_in <- element
        }
    }()

    go stage1(stage1_in, stage1_out)
    go stage2(stage1_out, nil)
}
```

Code example

Pipeline pattern in Go!

```
func main() {
    stage1_in := make(chan int)
    stage1_out := make(chan int)

    go func() {
        for _, element := range [4]int{1, 2, 3, 4} {
            stage1_in <- element
        }
    }()
}

go stage1(stage1_in, stage1_out)
go stage2(stage1_out, nil)
}
```

Code example

Pipeline pattern in Go!

```
func main() {
    stage1_in := make(chan int)
    stage1_out := make(chan int)

    go func() {
        for _, element := range [4]int{1, 2, 3, 4} {
            stage1_in <- element
        }
    }()
}

go stage1(stage1_in, stage1_out)
go stage2(stage1_out, nil)
```

Code example

Pipeline pattern in Go!

```
func stage1(in chan int, out chan int) {
    for {
        current := <- in
        op := current * 2
        out <- op
    }
}

func stage2(in chan int, out chan int) {
    for {
        current := <- in
        fmt.Println("Value: ", current)
    }
}
```

Code example

Pipeline pattern in Go!

```
func stage1(in chan int, out chan int) {
    for {
        current := <- in                                Read data from channel
        op := current * 2
        out <- op
    }
}

func stage2(in chan int, out chan int) {
    for {
        current := <- in
        fmt.Println("Value: ", current)
    }
}
```

Code example

Pipeline pattern in Go!

```
func stage1(in chan int, out chan int) {
    for {
        current := <- in
        op := current * 2          Perform operation
        out <- op
    }
}

func stage2(in chan int, out chan int) {
    for {
        current := <- in
        fmt.Println("Value: ", current)
    }
}
```

Code example

Pipeline pattern in Go!

```
func stage1(in chan int, out chan int) {
    for {
        current := <- in
        op := current * 2
        out <- op          Send data to next stage's channel
    }
}

func stage2(in chan int, out chan int) {
    for {
        current := <- in
        fmt.Println("Value: ", current)
    }
}
```

Code example

Pipeline pattern in Java...

Code example

Pipeline pattern in Java...

```
import java.util.ArrayList;
import java.util.List;

public interface Pipeline extends Stage {
    public void addStage ( Stage stage);

    public void addErrorStage (Stage stage);

    public void addFinalStage (Stage stage);
}

public interface PipelineContext {
    public List<Error> getErrors ();
}

public class PipelineContextAdaptor implements PipelineContext {

    private List<Error> m_errors = new ArrayList<Error> ();
    public List<Error> getErrors() {
        return m_errors;
    }

    public void addError (Error e){
        m_errors.add(e);
    }
}

public interface Stage {
    public void execute (PipelineContext context);
}

public class SequentialPipeline implements Pipeline {

    private List<Stage> m_stages = new ArrayList<Stage> ();
    private List<Stage> m_errorStages= new ArrayList<Stage> ();
    private List<Stage> m_finalStages= new ArrayList<Stage> ();

    public void addStage(Stage stage) {
        m_stages.add(stage);
    }

    public void addErrorStage(Stage stage) {
        m_errorStages.add(stage);
    }

    public void addFinalStage(Stage stage) {
        m_finalStages.add(stage);
    }

    public void execute(PipelineContext context) {
        for (Stage stage:m_stages){

            stage.execute(context);

            if (context.getErrors()!= null && !context.getErrors().isEmpty()){
                break;
            }
        }

        if (context.getErrors()!= null && !context.getErrors().isEmpty()){

            for (Stage errorStage: m_errorStages){
                errorStage.execute(context);
            }
        }

        for (Stage finalStage: m_finalStages){
            finalStage.execute(context);
        }
    }
}
```

Google Go

Can we use it on embedded platforms?

- We've seen that Go
 - allows for quicker development due to less boilerplate
 - simple syntax – few corner cases

Google Go

Can we use it on embedded platforms?

- We've seen that Go
 - is great for building parallel applications
 - channels as a language feature
 - go-routines (lightweight threads) as a language feature
 - compiler/runtime knows these features, **can find errors**

Google Go

Can we use it on embedded platforms?

- Go Concurrency
 - Straightforward implementation of the pipeline pattern
 - But other concurrency patterns work great too
 - Actors
 - Fork/Join
 - Master/Worker

Google Go

Can we use it on embedded platforms?

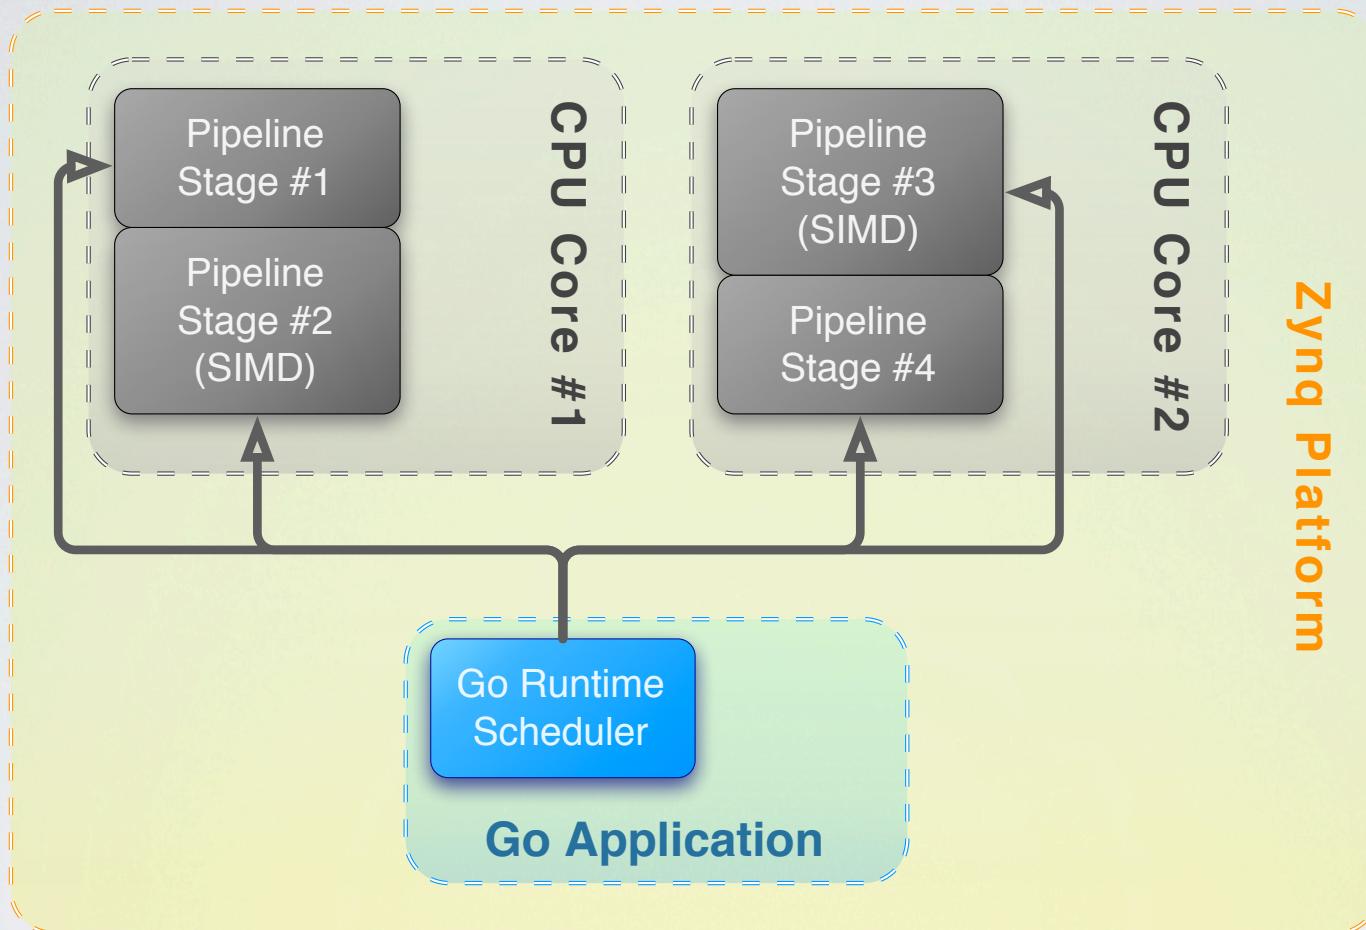
- But: simplicity comes at a cost
- Therefore, the question is
 - How big is the performance impact of using Go on embedded system?
 - And ultimately: Do the advantages of Go outweigh said performance impact?

Next Milestones

Next Milestones

- Technological unification on Linux
 - Determine if Go is a viable alternative to C
- Re-implementation of obstacle detection in Go or C
 - Pipeline Pattern and SIMD for throughput optimization

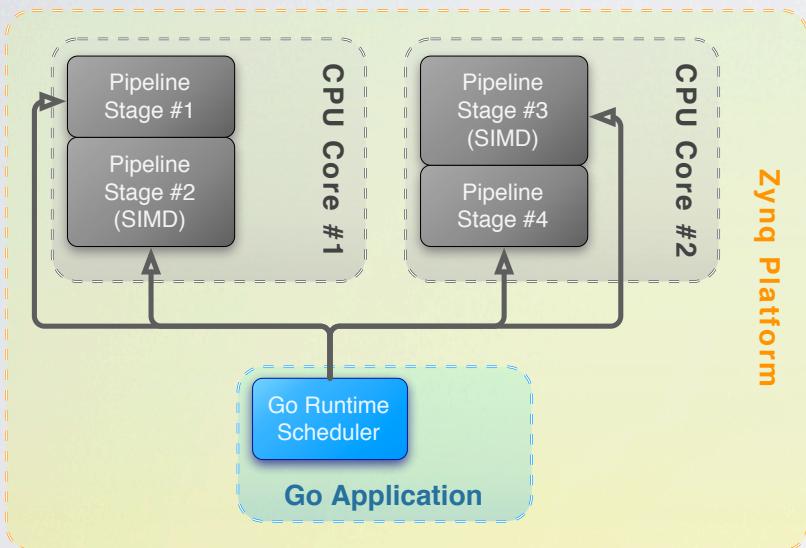
The Big Picture



Legend:

Go-Routine

The Big Picture



Legend:

Go-Routine

- Processing pipeline distributed to multiple cores
- Some stages use SIMD
- Go Runtime schedules the go-routines

The End
Any questions?

References

FAUST Project (2013). FAUST - Fahrerassistenz und autonome Systeme.

URL: <http://faust.informatik.haw-hamburg.de> (last accessed: 2013-05-07)

Kejariwal, A., Veidenbaum, A. V., Nicolau, A., Girkar, M., Tian, X. and Saito, H. (2009). On the exploitation of loop-level parallelism in embedded applications, *ACM Transactions on Embedded Computing Systems* 8(2): 1–34. ISSN: 15399087.

URL: <http://portal.acm.org/citation.cfm?doid=1457255.1457257>

Mattson, T. G., Sanders, B. A. and Massingill, B. L. (2010). *Patterns for Parallel Programming*, 6th edn, Addison-Wesley, Westford, Massachusetts. ISBN: 0-321-22811- 1.

Pike, R. (2012). Go at Google, *Proceedings of the 3rd annual conference on Systems, programming, and applications: software for humanity - SPLASH '12* p. 5.

URL: <http://dl.acm.org/citation.cfm?doid=2384716.2384720>

Preud'Homme, T., Sopena, J., Thomas, G. and Folliot, B. (2012). An Improvement of OpenMP Pipeline Parallelism with the BatchQueue Algorithm, *2012 IEEE 18th International Conference on Parallel and Distributed Systems* (i): 348–355.

URL: <http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber= 6413677>

Images

[1] URL:

<http://de.manu-systems.com/URG-04LX-UG01.shtml> (last accessed on 2013-06-02)

[2] URL:

http://www.wall321.com/Cars/Volkswagen/cars_vehicles_volkswagen_1280x800_wallpaper_44233/download_1920x1200 (last accessed on 2013-05-21)

[3] URL:

<http://www.wallpaperdev.com/wallpaper/1600x1200/cute-cat-by-ashish-8511.html> (last accessed on 2013-06-03)